

# NN DL Project

Group Members: Liyao Cai, Ying Liu, Modi Li

# Introduction

- Primary goal:
  - Classify the sentiment of IMDB movie reviews into positive or negative categories.
  - Explore and compare the performance of models in textual sentiment analysis
- Models:
  - RNN
  - LSTM
  - Transformer
- Dataset: IMDB dataset provided by the Keras library
  - 50,000 labeled movie reviews with positive and negative sentiments
- Why?
  - It provides a hands-on opportunity to understand and compare the capabilities of traditional and modern NLP models in handling sequential data.

# Outline

- Data preprocessing
- Model discussion
  - RNN
  - LSTM
  - Transformer
- Model results
- Comparative analysis

# Data Preprocessing

1. Define the max features and sequence length
2. Load and encode
3. Combine and split
4. Pad or truncate the merged data

```
max_features = 10000
max_len = 500 # Maximum length of each comment

# Load the data and convert the reviews into integer-encoded sequences
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_features)


# Combine all data
x_full_train = np.concatenate([x_train, x_test])
y_full_train = np.concatenate([y_train, y_test])

# The first 40,000 records are used as the new training set, and the remaining 10,000 records are used as the new test set.
x_train_new = x_full_train[:40000]
y_train_new = y_full_train[:40000]

x_test_new = x_full_train[40000:]
y_test_new = y_full_train[40000:]

# Pad or truncate the merged data
x_train_new = sequence.pad_sequences(x_train_new, maxlen=max_len)
x_test_new = sequence.pad_sequences(x_test_new, maxlen=max_len)
```

# Model: RNN



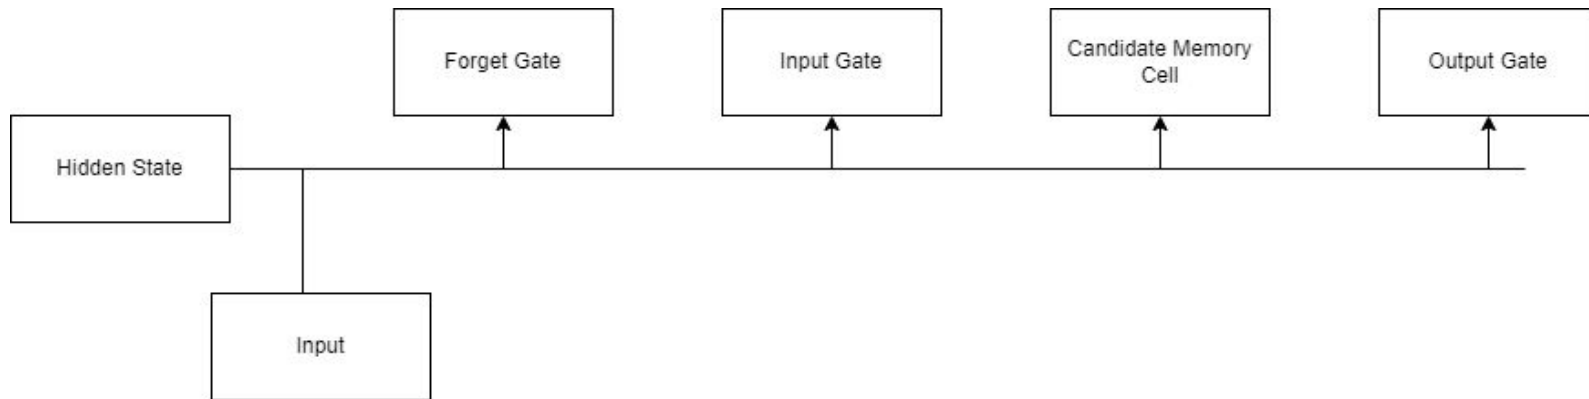
- **Recurrent Neural Networks (RNNs)** are a type of neural network designed for processing sequential data. RNNs have a feedback loop that allows information to persist across time steps. This makes them ideal for tasks involving sequences, such as:

- Text data
- Time-Series Data
- Speech and Audio

## Arrow Connections

- Input to Hidden State
- Hidden state to Hidden state
- Hidden state to Output

# Model: LSTM



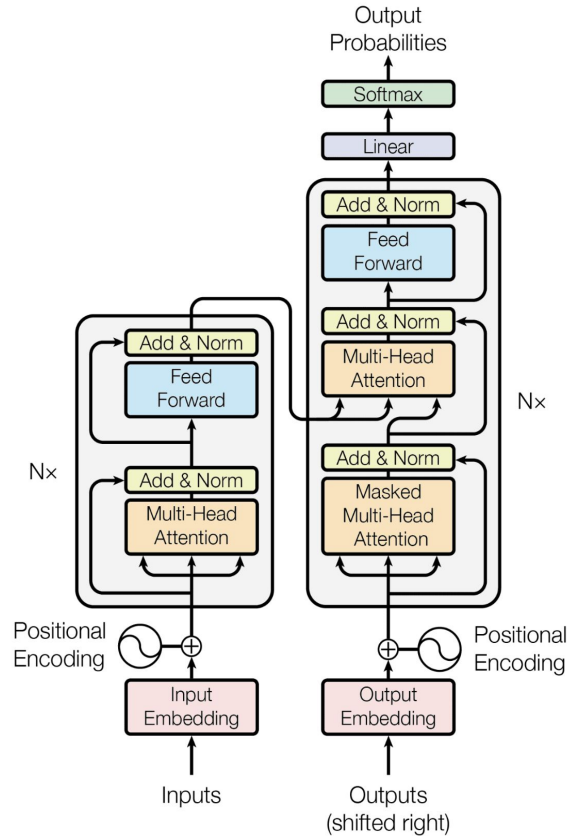
LSTM (Long Short-Term Memory) is a special type of RNN designed specifically to solve the long-term dependency problem of RNN.

Important parts

1. Memory Cell
2. Gating Mechanism
  - a. Forget Gate
  - b. Input Gate
  - c. Candidate Memory Cell
  - d. Output Gate

# Model: Transformer

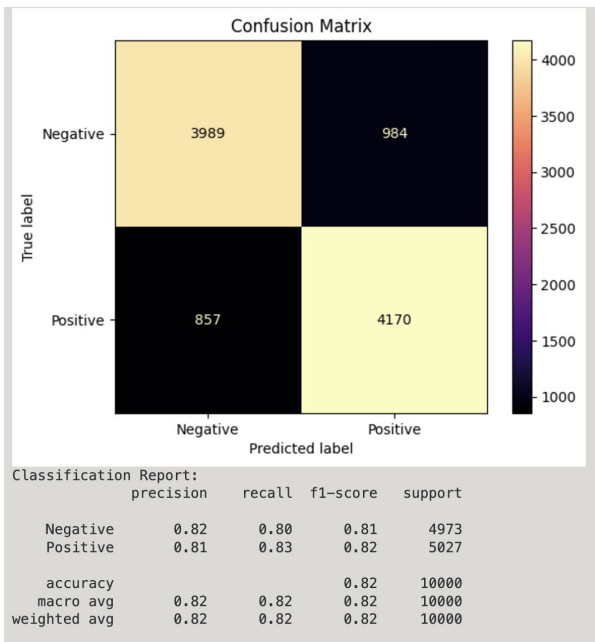
Transformer is a deep learning model architecture based on the Attention mechanism. It solves the shortcomings of traditional RNN/CNN in long sequence modeling and is widely used in Natural Language Processing, image processing and time series tasks.



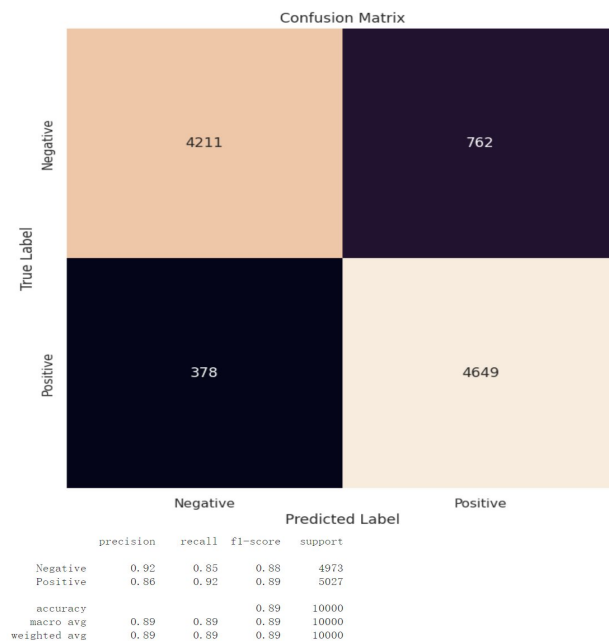
## Important parts

1. **Positional encoding:** Using sine and cosine positional encoding helps the model capture sequence order information.
2. **Multi-Head Attention:** Capturing the correlation between positions in the input sequence.


# Results



RNN



LSTM



Transformer



# Result Comparison

Model / Data Label	Precision	Recall	F1 Score	Accuracy	Running Time
RNN - Positive	0.82	0.80	0.81	0.82	6m
RNN - Negative	0.81	0.83	0.82		
LSTM - Positive	0.92	0.85	0.88	0.89	12m 44s
LSTM - Negative	0.86	0.92	0.89		
Transformer - Positive	0.89	0.89	0.89	0.89	2h 54m 26s
Transformer - Negative	0.90	0.89	0.89		

# Result Analysis

- RNN
  - Has lowest metrics
  - Shortest running time
  - Reason: struggles with long-range dependencies due to the vanishing gradient problem
  - Basic model with limited ability to capture context in texts
- LSTM
  - Has higher metrics than RNN
  - Longer running time than RNN
  - Very high recall for negative texts, showing better recognition of negative text patterns
  - Reason: overcomes RNN's limitations using the memory cells
  - Better at long-term dependencies
- Transformer
  - Has very high and most balanced performance among all models
  - Longest running time
  - Reason: uses the attention mechanism and captures long-range dependencies and the context more effectively while enabling parallel processing

# DistilBERT

- A small, fast, cheap and light Transformer model trained by distilling BERT base

Epoch 1/4

```
WARNING:tensorflow:AutoGraph could not transform <function infer_framework at 0x7f9f2e17a5f0> and will run it as-is.
Cause: for/else statement not yet supported
To silence this warning, decorate the function with @tf.autograph.experimental.do_not_convert
WARNING: AutoGraph could not transform <function infer_framework at 0x7f9f2e17a5f0> and will run it as-is.
Cause: for/else statement not yet supported
To silence this warning, decorate the function with @tf.autograph.experimental.do_not_convert
1250/1250 [=====] - 2304s 2s/step - loss: 0.7252 - accuracy: 0.5000 - val_loss: 0.6932 - val_accuracy: 0.5027
Epoch 2/4
1250/1250 [=====] - 2208s 2s/step - loss: 0.6935 - accuracy: 0.5038 - val_loss: 0.6938 - val_accuracy: 0.5027
Epoch 3/4
1250/1250 [=====] - 2228s 2s/step - loss: 0.6936 - accuracy: 0.4999 - val_loss: 0.6937 - val_accuracy: 0.5027
Epoch 4/4
1250/1250 [=====] - 2226s 2s/step - loss: 0.6935 - accuracy: 0.5034 - val_loss: 0.6932 - val_accuracy: 0.4973
```

```
313/313 [=====] - 182s 582ms/step - loss: 0.6932 - accuracy: 0.4973
```

Test Loss: 0.693188488483429

Test Accuracy: 0.49729999899864197

- Reasons
  - No find-tuning on the large dataset
  - Maybe the pre-trained model was not trained on sentiment specific data

Thank You